

Analytic Approximations for Real-Time Area Light Shading

Pascal Lecocq, Arthur Dufay, Gaël Sourimant, and Jean-Eudes Marvie

Abstract—We introduce analytic approximations for accurate real-time rendering of surfaces lit by non-occluded area light sources. Our solution leverages the *Irradiance Tensors* developed by Arvo for the shading of *Phong* surfaces lit by a polygonal light source. Using a reformulation of the 1D boundary edge integral, we develop a general framework for approximating and evaluating the integral in constant time using simple peak shape functions. To overcome the *Phong* restriction, we propose a low cost edge splitting strategy that accounts for the spherical warp introduced by the half vector parametrization. Thanks to this novel extension, we accurately approximate common microfacet BRDFs, providing a practical method producing specular stretches that closely match the ground truth in real-time. Finally, using the same approximation framework, we introduce support for spherical and disc area light sources, based on an original polygon spinning method supporting non-uniform scaling operations and horizon clipping. Implemented on a GPU, our method achieves real-time performances without any assumption on area light shape nor surface roughness.

Index Terms—Area light, shading, analytic, microfacet, axial moment, real-time

1 INTRODUCTION

ACCURATE real-time rendering of specular surfaces is a challenging task when considering area light source illumination. The difficulty resides in the evaluation of a two dimensional specular radiance integral for which no practical solution exists, except expensive Monte Carlo based sampling techniques. Most compelling solutions are found using *Most Representative Point (MRP)* approaches [1], [2], reducing the shading integration problem to a cheap point lighting calculation. However, these methods fail in preserving the specular highlight shape of underlying BRDFs and partial visibility above horizon is complicated to handle. Arvo [3] provides an exact analytic solution for the shading of glossy surfaces lit by a non-occluded polygonal light source. But its implementation relies on an expensive contour integration method, and is restricted to *Phong* surfaces. A recent and concurrent approach tackles this problem using *Linearly Transformed Cosine distributions (LTC)*[4]. However, the solution requires per-brdf precomputed tables built upon an expensive minimization technique.

Finding a solution combining accuracy, flexibility and real-time performances is a challenging problem with many expectations on high quality demanding applications such as lighting pre-viz tools, game engines or production renderers.

In this paper we address these shortcomings by leveraging the *Irradiance Tensors* developed by Arvo with accurate analytic approximations (Fig. 1). We further extend the

method to handle multiple axis-oriented *Cosine lobes* and overcome the *Phong* restriction, enabling support for microfacet BRDFs. Finally, we introduce an original polygon spinning method allowing surfaces shaded by spherical and disc area lights using the same mathematical framework.

Our contributions are:

- A general framework for approximating and evaluating the edge contour integrals in $O(1)$ time instead of $O(n)$ using simple and integrable peak shape functions.
- An analytical approximation for the multiple product of axis-oriented *Cosine lobes* that enables the integration of more complex BRDFs over spherical polygons.
- A low-cost edge splitting strategy for handling the warp distortion introduced by the half vector parameterization that enables microfacet BRDF support.
- An original spinning algorithm enabling spherical and disc area lighting and leveraging our approximations, that supports non-uniform scale operations.

2 RELATED WORK

Direct illumination from area light sources has been addressed in various ways in the last decades. We review in this section the related works we think most relevant to our approach, with a focus on the techniques addressing the integration of the specular term with real-time rendering constraints.

Monte Carlo Integration. Monte-Carlo integration techniques are a common approach to numerically compute complex integrals based on probabilistic sampling strategies. For direct area light illumination problems, samples are drawn either considering the solid angle sustained by the area shape [5], [6], considering importance sampling of

• The authors are with the Research & Innovation Department of Technicolor, 975 av des Champs-Blancs, Cesson-Sevigne 35576, France. E-mail: {pascal.lecocq, arthur.dufay, gael.sourimant, jean-eudes.marvie}@technicolor.com.

Manuscript received 25 Oct. 2016; revised 13 Jan. 2017; accepted 16 Jan. 2017. Date of publication 23 Jan. 2017; date of current version 29 Mar. 2017. Recommended for acceptance by K. Subr and L.-Y. Wei. For information on obtaining reprints of this article, please send e-mail to: reprints@ieee.org, and reference the Digital Object Identifier below. Digital Object Identifier no. 10.1109/TVCG.2017.2656889




Fig. 1. Our analytic approximations provide real-time performances to specular area lighting at a quality close to the ground truth (a, timing measured on a GPU NVIDIA GTX 980Ti). We support polygonal light source shapes of any kind (even animated ones) (b), and surfaces described by *Phong* or microfacet BRDF models (c). Our method easily extends to spherical and disc area lights with non uniform scaling (d, here with a microfacet distribution).

the surface BRDF, or using a combination of both to reduce the variance in presence of specular surfaces. Despite this sampling effort, these methods require a huge amount of samples to converge to a noise free result, hardly compatible with real-time rendering constraints.

Another common approach is to approximate area light sources using a set of *Virtual Point Lights* (VPLs) [7], reducing the specular term integration to a many point lights calculation. Clustering methods [8] have been proposed to further reduce the algorithm complexity of VPLs with successful application in real-time rendering [9], [10]. However, these solutions are usually restricted to low frequency illumination problems such as diffuse or weakly glossy surfaces to limit the sampling count and maintain good real-time performances. Rendering high frequency illumination with these methods is still a challenging problem only addressed using huge number of samples or expensive integration techniques far from real-time rendering considerations.

Most Representative Point. MRP approaches alleviate the costly sampling techniques by identifying a representative point on the area light that most contributes to the illumination. The method reduces the shading integration problem to a single point lighting calculation providing a practical solution for real-time rendering. Early works on the method can be found in [11] for *Phong* area lighting. The MRP here is defined as the closest point from the viewing reflection direction. Instead, Drobot [1] considers a point in the area of intersection between an area light and a cone with aperture parameterized by the surface roughness. Karis [2] addresses the problem of energy conservation and uses a modification of the specular distribution to better match intensity highlight of specular microfacet models. However, these approaches have several drawbacks. The highlight shape with a *Phong* BRDF is decently approximated, but becomes inaccurate when considering microfacet BRDFs. Horizon handling is yet another issue. The MRP approximation works well with simple geometric light emitters but calculation get more complex when the light source is clipped above the horizon plane.

Analytic Approaches. Other approaches try to derive an exact analytic solution of the shading integral, or at least a decent approximation. Bao and Peng [12] approximate the double integral with 2D polynomials using a low degree Taylor series expansion, limiting their method to low

exponent *Phong* surfaces. Tanaka and Takahashi [13] extend the linear area light method of Poulin [14] and decompose the solid angle into 1D signed integrals along edge great circle. Each 1D integral is then evaluated using a Chebyshev polynomial approximation, restricting the method to low frequency *Phong* surfaces. The *Irradiance Tensors* developed by Arvo [3] provide an exact analytic solution for the direct illumination of glossy surfaces lit by a polygonal light source. Using tensor theory and Stokes contour integration, the shading integral is decomposed into a sum of signed 1D integrals along the spherical boundary edges of the polygonal light. Each edge integral is then evaluated analytically using a linear time algorithm bound to the *Phong* shininess n . A practical implementation for real-time graphics, including horizon clipping, can be found in [15]. Despite its accuracy, the method only works for *Phong* surfaces and its usage in real-time rendering applications is limited to weakly glossy surfaces due its $O(n)$ time complexity.

Spherical Gaussians (SGs). SGs are spherical functions used in many lighting problems such as *environment lighting* or *global illumination* with subsequent derivations for real-time area light illumination. Wang et al. [16] approximate a spherical area light with an SG providing a closed-form expression for the integral product with an SG-approximated BRDF. To handle microfacet BRDFs, the spherical warp introduced by the half vector transform is approximated using a single isotropic SG. However this method fails to represent the elongated specular stretches at grazing angles. Xu et al. [17] approximate the spherical warping using *Anisotropic Spherical Gaussians* (ASGs). A practical implementation for spherical light source illumination can be found in [18]. These methods have two main limitations. First, the spherical warp approximation supposes an isotropic light source emitter. Second, highly glossy surfaces tend to reveal a Gaussian shape due to the area light approximation as an SG. Close to our approach, Wang et al. [19] use a piece-wise linear approximation of the SG for polygonal visibility evaluation, reducing the 1D edge integrals to analytic expressions. Conversely, Xu et al. [20] use an edge parameterization on the parallel plane to derive 1D edge integral expressions evaluated using a piece-wise linear approximation. Both method require a piece-wise decomposition of the integration domain. Furthermore, they are restricted to isotropic SGs only and fail to represent the anisotropy of microfacet distributions.

Linear Transformed Cosine. A recent and concurrent approach proposed by Heitz et al. [4] uses linear transforms of a clamped Cosine distribution (referred hereafter as *LTC*) to approximate isotropic BRDFs including microfacets. As a result, by applying the inverse transform on the polygon, the shading operation is reduced to an analytical form factor calculation. *LTC* is accurate and simpler to evaluate than ours. However, matching a BRDF requires the pre-computation and storage of the transformation matrices for the set of incident directions and roughness values. The need for storage and memory access is not always desirable especially on low-end GPUs. Our solution is fully analytic and does not require any pre-computation step or storage. Furthermore, our method enables integration in the half-vector space, providing higher degree of liberty. In that sense, both methods are complementary.

3 OUR APPROACH

Our method builds upon *Irradiance Tensors* and the contour integration method developed by Arvo [3], that we briefly recall in Section 4. This approach represents several challenges.

The first challenge is to get around the $O(n)$ time bottleneck for real-time rendering efficiency. We tackle this problem by rewriting the 1D integrals in a more concise way (Section 5) allowing to settle for an accurate $O(1)$ time approximation using a rational peak shape integration framework (Section 6). Unlike *Chebyshev* or *Fourier* approximations, our approach is bound to only 1 or 2 rational functions and doesn't suffer from any ringing artifacts.

Second, the integration of more complex distributions over a spherical polygon combining several axis-oriented *Cosine lobes* is yet another challenge. By borrowing operators from *Spherical Gaussians*, we derive simple analytic expressions (Section 7) approximating accurately the multiple product of *Cosine lobes*.

Finally, the last challenge is to overcome the *Phong* BRDF restriction and give support for more plausible BRDFs. The half-vector parameterization found in microfacet theory introduces a spherical distortion which can be difficult to predict using non isotropic polygonal light sources. Based on observations from great circle distortions, we can faithfully approximate this spherical warp using a polygonal approach (Section 8) in a more flexible way than previous methods.

4 IRRADIANCE TENSORS AND THE EDGE INTEGRAL

The *Irradiance Tensors* developed by Arvo [21] provide a useful framework for the analytic integration of polynomials over the sphere S^2 . These polynomials correspond to n th order monomial expressions described by an axis-oriented cosine lobe distribution. The integration of this expression over a spherical region $\Omega_A \subset S^2$ yields to the definition of n th order *axial moment* about an \mathbf{r} axis

$$M^n(\Omega_A, \mathbf{r}) = \int_{\Omega_A} (\mathbf{u} \cdot \mathbf{r})^n d\mathbf{u}. \quad (1)$$




Fig. 2. The integral of a cosine lobe distribution axis, oriented toward \mathbf{r} , over a spherical region Ω_A is reduced to a 1D contour integration over the boundary of Ω_A . A closed-form expression is given for polygonal light source using a parameterization of the i th spherical edge in the local base $\mathbf{v}_i, \mathbf{t}_i$ (the red dotted line depicts here the great circle supporting the edge $(\mathbf{v}_i, \mathbf{v}_{i+1})$).

Using tensors product and Stokes theorem, Arvo developed the axial moment expression as a 1D contour integration over the projected area light boundary Ω_A . Considering a polygonal light source, a closed-form expression for the 1D integrals can be obtained following a parameterization of spherical edges along great circles (see Fig. 2). Let consider a polygon with m boundary edges. Following the notations depicted in Fig. 2, the closed-form expression is given as follows

$$(n+1)M^n(\Omega_A, \mathbf{r}) = z_n \Omega_A - \sum_{i=0}^m (\mathbf{n}_i \cdot \mathbf{r}) F(\Phi_i, c_i, \delta_i, n-1), \quad (2)$$

with

$$c_i = \sqrt{a_i^2 + b_i^2}; \quad \delta_i = \tan^{-1}(b_i/a_i)$$

$$a_i = \mathbf{v}_i \cdot \mathbf{r}; \quad b_i = \mathbf{t}_i \cdot \mathbf{r},$$

and

$$F(\Phi, c, \delta, n) = \sum_{k=0}^{\frac{n-1+z_n}{2}} c^{2k+1-z_n} \int_{-\delta}^{\Phi-\delta} (\cos \phi)^{2k+1-z_n} d\phi, \quad (3)$$

with $z_n = 1 - (n \bmod 2)$.

For a complete description of *Irradiance Tensors* and how to come to this expression, the reader shall refer to [21] and [3]. Note also that we use a slightly different notation compared to Arvo to ease the mathematical derivations further developed in the next sections.

The sum of 1D integrals in F is evaluated in closed-form using a recurrence algorithm of complexity $O(n)$ time per edge, n being the *Phong* exponent. Implemented on a GPU, the method works well for weakly glossy surfaces ($n < 40$) but the performance drops as the *Phong* exponent increases, and becomes impractical for highly glossy surface ($n > 1,000$). To reduce the evaluation cost for high *Phong* exponents, Arvo suggested early termination of the iteration loop once a desired relative accuracy is reached. But, from our experience, we observe severe performance drop-off, especially at grazing view angles of the surface because of a high number of iterations necessary to reach

TABLE 1
Notations Used Throughout This Paper

Symbol	Description
$L(x, \mathbf{v})$	radiance scattered from point x toward direction \mathbf{v}
Ω_A	region solid angle sustained by area light source A
\mathbf{n}	surface normal
\mathbf{v}	normalized view vector
\mathbf{r}	normalized reflected view vector
\mathbf{h}	normalized halfway vector given by $(\mathbf{i} + \mathbf{v})/ \mathbf{i} + \mathbf{v} $
m	number of boundary edges on the polygonal light
\mathbf{v}_i	spherical projection of the i th vertex of the polygonal light
n	cosine lobe exponent
\mathbf{n}_i	edge outer normal given by $(\mathbf{v}_i \times \mathbf{v}_{i+1})/ \mathbf{v}_i \times \mathbf{v}_{i+1} $
\mathbf{t}_i	edge tangent vector given by $\mathbf{t}_i = \mathbf{v}_i \times \mathbf{n}_i$
Φ_i	edge arc length
F	spherical edge integral
\tilde{F}	approximated spherical edge line integral

the desired accuracy. The difficulty to predict the performance makes it a nonviable solution for real-time rendering considerations. In a practical GPU implementation, the edge integral F should be ideally evaluated in $O(1)$ whatever the Phong exponent.

5 REFORMULATION OF THE EDGE INTEGRAL

We propose to replace the costly edge integral evaluation of Equation (3) by a cheap and accurate analytic approximation that allows constant time evaluation with any Phong shininess n . Setting an accurate approximation requires the knowledge or at least the intrinsic characteristics of the integrand function. A common approach is to probe the edge integrand to extract these characteristics. However, in its present form, this requires the evaluation of the sum of the integrand terms. We propose to rewrite the edge integral in a different form in order to get a simpler and more compact expression. To that end, we first introduce a term f and a temporary term q defined as follow

$$f(\phi, c, n) = \begin{cases} c \cos \phi & q = \begin{cases} (n-1)/2 & \text{if } n \text{ is odd} \\ n/2 & \text{if } n \text{ is even} \end{cases} \end{cases}$$

By switching the sum and integral operators and by using the terms introduced above, the edge integral F defined in Equation (3) can be re-written as follow

$$F(\Phi, c, \delta, n) = \int_{-\delta}^{\Phi-\delta} f(\phi, c, n) \sum_{k=0}^q (c \cos \phi)^{2k} d\phi. \quad (4)$$

The sum exhibits a geometric series of the form x^{2k} with a generic formula

$$\sum_{k=0}^q x^{2k} = (x^{2(q+1)} - 1)/(x^2 - 1).$$

This allows us to cancel out the sum and get a single function to integrate after substituting the temporary variable q .

$$F(\Phi, c, \delta, n) = \int_{-\delta}^{\Phi-\delta} \frac{(c \cos \phi)^{n+2} - f(\phi, c, n)}{(c \cos \phi)^2 - 1} d\phi. \quad (5)$$




Fig. 3. The edge integrand I reveals peak shape functions of various height and width. The core idea is to approximate I using a simple and integrable peak shape function with same characteristics.

This reformulation allows the evaluation of the integrand term in constant time. Another advantage is that it enables smooth representation of non n integer values, especially for low n exponents. Though no indefinite integral exists, an accurate analytic approximation can be obtained from our reformulation.

6 ACCURATE ANALYTIC APPROXIMATIONS

Let us consider the integrand term from the edge integral in Equation (5)

$$I(\Phi, c, n) = \frac{(c \cos \phi)^{n+2} - f(\phi, c, n)}{(c \cos \phi)^2 - 1}. \quad (6)$$

According to Fig. 3, we observe that the shape of I corresponds to symmetric peak shape functions of various height and width depending on parameters c, n having a minimum reached at $\phi = \pm\pi/2$ and a maximum at $\phi = 0$.

The core idea of our method is to approximate I using peak shape functions described by simple rational expressions with known analytic integration. The approximation relies on a simple fitting procedure that maps a peak shape function to the integrand I characteristics such as the minimum, maximum and width.

Following Equation (6), a closed formulation is given for the minimum and maximum values

$$I_{\min}(c, n) = \begin{cases} 0 & n \text{ odd} \\ 1 & n \text{ even} \end{cases} \quad I_{\max}(c, n) = \begin{cases} \frac{c^{n+2}-c}{c^2-1} & n \text{ odd} \\ \frac{c^{n+2}-1}{c^2-1} & n \text{ even} \end{cases}$$

Half Width Estimation. The width is defined as the Half Width at Half Maximum (HWHM) which corresponds to the abscissa x_w such as

$$I(x_w, c, n) = (I_{\max} - I_{\min})/2 - I_{\min}.$$




Fig. 4. Approximation of the half width x_w of integrand I for various values of n odd (left) and n even (right).

However, finding a closed-form expression for x_w is somewhat more difficult. One approach could consist in storing pre-computed values for x_w in a 2D table for discrete entries (c, n) . Another approach is to settle an analytic approximation. Following experimental measurement studies, we found that x_w can be empirically approximated as follow

$$x_w(c, n) \approx \begin{cases} \frac{\pi}{3} \sqrt{1 - (c - \frac{c}{n})^2} & n \text{ odd} \\ \frac{\pi}{4} \left(1 - (c - \frac{c}{n-1})^{2.5}\right)^{0.45} & n \text{ even} \end{cases} \quad (7)$$

Even if this is a rough estimation (Fig. 4), the fitting procedure, described in next section, will guarantee that our approximation will pass through the point $(x_w, I(x_w, c, n))$.

6.1 General Integration Framework

We derive a general framework for approximating and evaluating Equation (3) by means of generic peak shape functions. To that end, we first consider a generic peak function P , defined by a minimum P_{\min} , a maximum P_{\max} and width P_w . An accurate approximation of I can be obtained by adjusting P to the same characteristics of I . The fitting procedure consists in a scaling, offsetting and width adjustment defined as follow

$$\tilde{I}(\phi, c, n) = \frac{I_{\max} - I_{\min}}{P_{\max} - P_{\min}} (P(\phi, x_w) - P_{\min}) + I_{\min}.$$

Note that function parameters have been omitted for brevity. We can further reduce this expression by packing all the constant terms together

$$\tilde{I}(\phi, c, n) = s P(\phi, x_w) + t, \quad (8)$$

with $s = (I_{\max} - I_{\min}) / (P_{\max} - P_{\min})$ and $t = I_{\min} - s P_{\min}$.

A general solution for the evaluation of Equation (3) is then given by

$$\tilde{F}(\Phi, c, \delta, n) = s \int_{-\delta}^{\Phi-\delta} P(\phi, x_w) d\phi + t\Phi. \quad (9)$$

6.2 Peak-Shape Functions Approximation

We studied several peak shape function families P with indefinite integrals simple enough to avoid time-consuming evaluation and providing an accurate estimate of Equation (6). We validated the accuracy of our approximations with ground truth comparison by implementing an energy-conserving

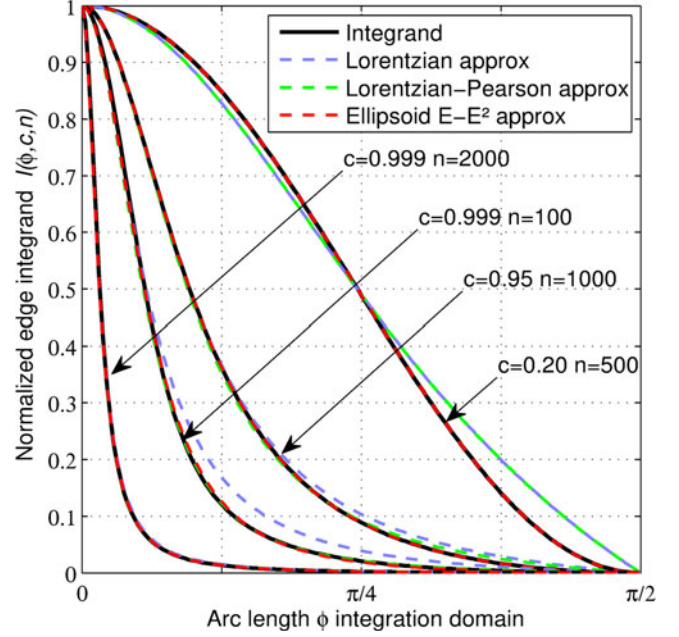


Fig. 5. Approximation of the edge integrand I using peak shape functions for various values of c and n . The *Lorentzian* function approximates edge integrands fairly well but lacks accuracy in the tail of I . The *Lorentzian-Pearson* better approximates I , but it remains inaccurate for large width values. The *ellipsoid* approximation provides the best accuracy whatever the width of the function.

single-axis *Phong* model using a single-axial moment evaluation expressed as follow

$$L(x, \mathbf{v}) = \int_{\Omega_A} f_{\text{Phong}}(\mathbf{i}, \mathbf{v})(\mathbf{n} \cdot \mathbf{i}) d\mathbf{i} = \rho_s \frac{n+1}{2\pi} M^n(\Omega_A, \mathbf{r}). \quad (10)$$

Horizon Clipping. Horizon clipping takes into account the energy loss when the area light is partially below the horizon. While the clipping procedure was not explicitly addressed by [3], a practicable implementation can be found in [15]. We adopt the same procedure in our implementations.

6.2.1 Lorentzian Approximation

The simplest approximation can be found by means of a *Lorentzian* peak shape function

$$P(\phi, c, n) = \frac{1}{1 + a\phi^2} \quad \text{with} \quad \int P = \frac{1}{\sqrt{a}} \tan^{-1}(\sqrt{a}\phi). \quad (11)$$

We use Equation (7) to compute the fitting point $I(x_w, c, n)$ that roughly corresponds to the half maximum of I . Solving the equation $I(x_w, c, n) = \tilde{I}(x_w, c, n)$ yields to resolution of unknown parameter a .

$$a = \frac{1 - y_w - \frac{4x_w^2}{\pi^2}}{y_w x_w^2} \quad \text{with} \quad y_w = \frac{I(x_w) - I_{\min}}{I_{\max} - I_{\min}}.$$

By replacing the integral term in Equation (9) by the one defined in Equation (11), we obtain an analytic approximation for F evaluated in constant-time

$$\tilde{F} = \frac{s}{\sqrt{a}} (\tan^{-1}(\sqrt{a}(\Phi - \delta)) - \tan^{-1}(-\delta\sqrt{a})) + t\Phi.$$




Fig. 6. Comparison of our approximations against ground truth reference images. The right column depicts a scaled version of the images on left one, so as to make differences visible. While light leaks are clearly visible for the Lorentzian and slightly visible for the Lorentzian-Pearson, they are hardly noticeable on the scaled version for the ellipsoid approximation.

Noting that $\tan^{-1}x - \tan^{-1}y = \tan^{-1}\left(\frac{x-y}{1+xy}\right) \bmod \pi$, this expression can be further reduced into a single arctangent evaluation to save GPU instructions.

$$\tilde{F} = \frac{s}{\sqrt{a}} \tan^{-1}\left(\frac{\sqrt{a}\Phi}{1+a(-\delta)(\Phi-\delta)}\right) \bmod \pi + t \Phi. \quad (12)$$

Error Analysis. The Fig. 6 shows that the Lorentzian approximation is fairly accurate and close to the ground truth whatever the roughness of the surface. However, we can observe slight light leaks around the highlight shape most noticeable when increasing the overall intensity. A careful observation of the Lorentzian approximation plots in Fig. 5 shows that the error results from an overestimation of the function I around the tail.

6.2.2 Lorentzian-Pearson VII Approximation

A better approximation around the tail can be found by combining the *Lorentzian* function with a second peak function with a shorter tail. The idea is to encompass the integrand function in the tail area with these two approximations and find a blending factor from values picked in the tail. The second peak function is defined by a *Pearson VII* function corresponding to a *Lorentzian* function raised to a power m . We chose $m = 2$ which has an indefinite integral simple enough to avoid time-consuming computations

$$P = \left(\frac{1}{1+b\phi^2}\right)^2 \text{ and } \int P = \frac{\phi}{2(b+\phi^2)} + \frac{\tan^{-1}\left(\frac{\phi}{\sqrt{b}}\right)}{2\sqrt{b}}.$$

The *Pearson VII* function behaves exactly like a *Lorentzian* function on its superior part and has a shorter tail in its bottom part. However, finding the b parameter, such as $\tilde{I}(x_w) = I(x_w)$ requires the resolution a polynomial equation

of degree four involving complex computations. Fortunately, it turns out that the computation of b can be greatly simplified and save GPU computation time by reusing the a parameter computed for the *Lorentzian* approximation. From our experiments, we found that $b \approx a/2$ always enclosed the target integrand function I .

Linear Blending. Adjusting I_P to the same width than I gives us another approximation that underestimates I in the tail while preserving the fitting above it. The best approximation then sits between the two functions and can be found using a simple linear blending operation.

$$\tilde{I}_{LP}(\phi) = \alpha \tilde{I}_L(\phi) + (1 - \alpha) \tilde{I}_P(\phi), \quad (13)$$

where

$$\alpha = \frac{\tilde{I}_P(x_{\text{tail}}) - I(x_{\text{tail}})}{\tilde{I}_P(x_{\text{tail}}) - \tilde{I}_L(x_{\text{tail}})}.$$

The linear blend operation requires the evaluation of the integrand function I at a position x_{tail} located in the tail of the function. However, finding a closed-form expression for x_{tail} represents the same difficulty as for the *half width* estimation. Again, we use instead an empirical approximation

$$x_{\text{tail}} \approx x_w + 0.3946 x_w(0) \left(1 - (1 - x_w/x_w(0))^{12}\right). \quad (14)$$

Approximation Accuracy. The *Lorentzian-Pearson* approximation greatly improves the overall accuracy of the specular highlight and suppresses most observable artifacts. Although, we still experience subtle light-leaks on areas located outside the specular highlights as shown in Fig. 6. These leaks are occurring when the peak shape I is very large, i.e., when the value c is small. A closer look at the plot in Fig. 5 shows that the approximation is overestimated at integration domain bounds. Especially, at $\phi = \pm\pi/2$, the first derivative is null while the *Lorentzian-Pearson* approximation is not.

6.2.3 Ellipsoid Approximation

A better accuracy, especially at domain bounds, can be obtained using ellipsoid-based peak shape functions. These functions have the interesting property to behave like a *Lorentzian* but having a null first derivative at $\phi = \pm\pi/2$.

$$\text{Ellipsoid : } P_E = \frac{a}{1 + (a-1) \cos^2(\phi)} \quad (15)$$

$$\text{Indefinite integral : } \int P_E = \sqrt{a} \tan^{-1}\left(\frac{\tan \phi}{\sqrt{a}}\right). \quad (16)$$

Square Ellipsoid function

$$P_{E^2} = \left(\frac{b}{1 + (b-1) \cos^2(\phi)}\right)^2; \\ \int P_{E^2} = \frac{\sqrt{b}}{2} (b+1) \tan^{-1}\left(\frac{\tan \phi}{\sqrt{b}}\right) - \frac{b}{2c} (b-1) \sin(2\phi). \quad (17)$$

We follow exactly the same procedure described in Sections 6.2.1 and 6.2.2 to fit I_E and I_{E^2} to I and find the

TABLE 2
Rendering Times in Milliseconds per Edge on a GPU NVIDIA GTX 580 and GTX 980Ti with Rendering Accuracy for the Three Peak Shape Approximations

Method	Exponent	Time/edge (ms)		RMSE
Arvo (exact)	$n = 100$	13.6	0.41	n/a
	$n = 500$	49	1.57	n/a
	$n = 5,000$	476	9.8	n/a
Lor approx	$n = 100$	0.25	0.12	0.004354
	$n = 500$	0.25	0.12	0.005506
	$n = 5,000$	0.25	0.12	0.004128
Lor-Pear approx	$n = 100$	0.40	0.125	0.003641
	$n = 500$	0.40	0.125	0.003094
	$n = 5,000$	0.40	0.125	0.002551
Ellipsoid approx	$n = 100$	0.47	0.127	0.001500
	$n = 500$	0.47	0.127	0.001652
	$n = 5,000$	0.47	0.127	0.001014

best approximation using a linear blending. The parameter a for the first approximation I_E corresponds to

$$a = \frac{y_w(1 - \cos(\phi)^2)}{\cos(\phi)^2(1 - y_w)}. \quad (18)$$

For I_{E^2} , parameter b roughly follows

$$b \approx a \left(2.1 + 1.28 \frac{x_w}{x_w(0)} \right).$$

Approximation Accuracy. The *ellipsoid* approximation provides the best accuracy whatever the *width* of the function with unnoticeable artifacts as illustrated in Fig. 6.

6.3 Performance versus Accuracy Analysis

We implemented and tested our approximations on a GPU NVIDIA GTX 580. The Table 2 provides the rendering times in milliseconds per edge along with rendering accuracy measurements using a normalized RMSE. Measurements were done considering the processing of all screen pixels, representing the most critical case, at a 720p resolution. Note that the timings also include the double horizon clipping around \mathbf{n} and around \mathbf{r} .

As expected, the rendering time obtained with Arvo's solution increases with the exponent n , while remaining constant with our approximations. The *Lorentzian* approximation achieves the best performance while the *ellipsoid* is the most accurate with unnoticeable difference with the ground truth and a small computational overhead introduced by a GPU time-consuming tangent evaluation. The *Lorentzian* approximation can be sufficient most of the time for high performance demanding application such as games. For high quality demanding applications such as lighting pre-viz for production rendering, the *Lorentzian-Pearson* or the *ellipsoid* approximation are the best choices.

7 MULTIPLE-AXES MOMENTS EVALUATION

The *Irradiance tensors* allow evaluation of multiple axes using decomposition of tensor product. Arvo [3] proposed a closed form expression for the double-axis moment

described by the product of two cosine lobes of order n and 1. However, the generalization for arbitrary orders combining several axes is quite difficult, involving complex mathematical expressions and requiring per-axis moment evaluation.

To reduce the mathematical complexity of the problem, we demonstrate that the product of *Cosine lobes* can be accurately approximated by a single *Cosine lobe*, providing a practicable and efficient way for multiple-axes moment evaluation.

$$\int_{\Omega_A} (\mathbf{u} \cdot \mathbf{r}_1)^{n_1} (\mathbf{u} \cdot \mathbf{r}_2)^{n_2} \dots (\mathbf{u} \cdot \mathbf{r}_i)^{n_i} \approx \int_{\Omega_A} \mu (\mathbf{u} \cdot \mathbf{r})^n. \quad (19)$$

Such mathematical reduction has been first addressed by Meunier et al. [22]. Their approach uses a time-consuming L_2 minimization technique to determine the parameters (μ, r, n) approximating the product of two *Cosine lobes*. The results are pre-computed and stored in a table for a large collection of $(n_1, n_2, \angle(\mathbf{r}_1, \mathbf{r}_2))$ samples. Dealing with tables is not always desirable in terms of memory occupation and cache efficiency, especially on low-end GPU mobile devices. To avoid this problem, we propose a simple analytic approach which is fast and that does not require any pre-computations step or memory storage.

7.1 Product of Cosine Lobes Approximation

Our approach is built from the simple observation that *Cosine lobes* and *Spherical Gaussians* share many similarities in terms of shape and convolution operators. SGs are a useful mathematical tool for approximating many lighting problems. As a brief recall, an SG is a spherical function with the following form

$$G(\mathbf{u}, \mathbf{r}, \lambda, \mu) = \mu e^{\lambda(\mathbf{u} \cdot \mathbf{r} - 1)}.$$

An interesting property of SGs is that the product of two SGs is an another SG, computed exactly using simple analytic formulas. For a consistent notation with SG, let us introduce the spherical function C to represent a *Cosine lobe* with a magnitude μ .

$$C(\mathbf{u}, \mathbf{r}, n, \mu) = \mu (\mathbf{u} \cdot \mathbf{r})^n. \quad (20)$$

It turns out that a *Cosine lobe* can be fairly well approximated by an SG in most situations.

$$C(\mathbf{u}, \mathbf{r}, n, \mu) \approx G(\mathbf{u}, \mathbf{r}, \lambda, \mu). \quad (21)$$

We have also observed that the product of two *Cosine lobes* closely behaves like the product of two SGs due to the shape similarities. Our idea is to borrow product operators from SGs to derive a single *Cosine lobe* approximation from the product of two *Cosine lobes*. Our method consists in mapping an SG on each *Cosine lobe* and evaluate the parameters of the product in the SG domain. Then, we back-transform the results in the *Cosine lobe* domain by mapping a *Cosine lobe* on the resulting SG.

The mapping of a *Cosine lobe* from / to an SG is achieved by solving the equations such as C and G have the same width at half maximum. These parameters are computed exactly as follows

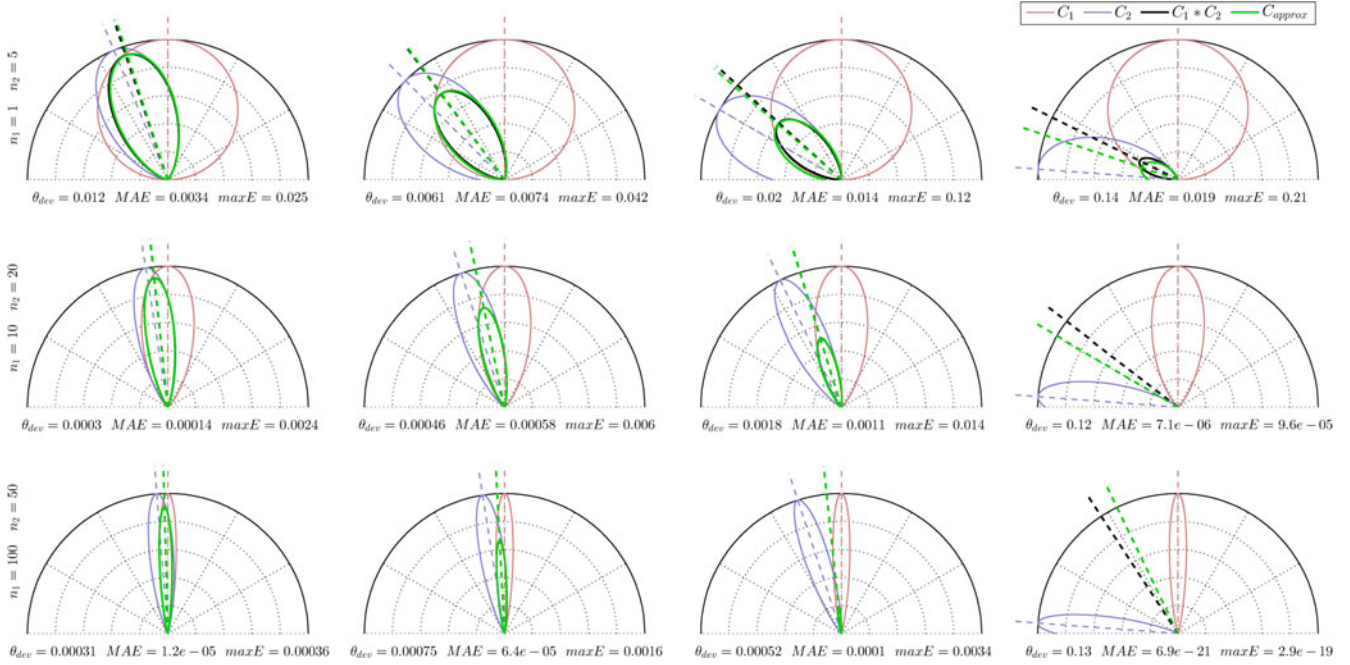


Fig. 7. Polar plots of the product (in black) of two *cosine lobes* (in red and blue) against our analytic approximation (in green) following the formulas given in Section 7.1 for various eccentricity and angle configurations. Our single lobe approximation closely matches the product of two lobes in most situations. The critical scenario arises when the two lobes are far apart. In this case, we observe a misalignment of the two lobes but the error remains un-noticeable due to the low magnitude of the product as illustrated in the rendered pictures.

$$\lambda = \frac{-\ln 2}{\sqrt{2}-1}; \quad n = \frac{-\ln 2}{\ln \frac{-\ln 2 + \lambda}{\lambda}}. \quad (22)$$

The product of two *Cosine lobes* C_1 and C_2 is then approximated as follows

$$\begin{aligned} C_1(\mathbf{r}_1, n_1, \mu_1) C_2(\mathbf{r}_2, n_2, \mu_2) &\approx G_1(\mathbf{r}_1, \lambda_1, \mu_1) G_2(\mathbf{r}_2, \lambda_2, \mu_2) \\ &\approx G(\mathbf{r}_p, \lambda_p, \mu_p) \\ &\approx C(\mathbf{r}_p, n_p, \mu_p), \end{aligned}$$

where

$$\begin{aligned} \mathbf{r}_p &= \frac{\mathbf{p}_m}{\|\mathbf{p}_m\|}; \quad \mathbf{p}_m = \frac{\lambda_1 \mathbf{r}_1 + \lambda_2 \mathbf{r}_2}{\lambda_1 + \lambda_2} \\ \lambda_p &= (\lambda_1 + \lambda_2) \|\mathbf{p}_m\| \\ \mu_p &= C_1(\mathbf{r}_p, \mathbf{r}_1, n_1, \mu_1) C_2(\mathbf{r}_p, \mathbf{r}_2, n_2, \mu_2). \end{aligned}$$

Note that we do not evaluate explicitly the SGs but only borrow their product operators for deriving our approximation. The only difference lies in the computation of *magnitude* μ_p computed exactly as the product of the two *Cosine lobes* at \mathbf{r}_p . Note that the same reasoning could be employed to approximate the product of anisotropic *Cosine lobes* by borrowing operators derived in [17]. We left this derivation for future works.

7.2 Results and Error Analysis

We compared our lobe product approximation C against an exact product of C_1 and C_2 for various exponents and angles in polar coordinates (see Fig. 7). In most situations, our single *Cosine lobe* approximation closely matches the product of two *Cosine lobes*. However, when the angle between the two lobes is very large, we observe a misalignment between the theoretical product and our approximation. Although,

in this situation, the error is balanced by the very low magnitude of the product, close to 0 except when the two lobes have a low exponent.

Double-Axis Phong Implementation. We further validated our approach by implementing the energy-conserving double-axis *Phong* model. According to the formulas given in Section 7.1, the model can be approximated by a single axial moment as follow

$$\begin{aligned} L(x, \mathbf{v}) &= \rho_s \frac{n+2}{2\pi} \int_{\Omega_A} (\mathbf{i}, \mathbf{r})^n (\mathbf{n} \cdot \mathbf{i}) d\mathbf{i} \\ &= \rho_s \frac{n+2}{2\pi} \int_{\Omega_A} C_1(\mathbf{i}, \mathbf{r}, n, 1) C_2(\mathbf{i}, \mathbf{n}, 1, 1) d\mathbf{i} \\ &\approx \rho_s \frac{n_p+1}{2\pi} \int_{\Omega_A} C(\mathbf{i}, \mathbf{r}_p, n_p, \mu_p) d\mathbf{i} \\ &\approx \rho_s \mu_p \frac{n_p+1}{2\pi} M^{n_p}(\Omega_A, \mathbf{r}_p). \end{aligned}$$

This approximation is energy-conserving thanks to the new normalization factor resulting from the *Cosine* product parameterization and the axial moment normalization.

Images in Fig. 8 show a rendering comparison of the double-axis *Phong* approximation against a ground truth solution. The visual difference is unnoticeable whatever the surface *Phong* exponent or the angular configuration. As predicted, the error occurring when the angle of the two lobes is large is visually dismissed by the low magnitude of the product.

8 EXTENSION TO MICROFACET BRDFs

The limitation to the *Phong* specular BRDF is a hard constraint for *Irradiance Tensors*. Most production renderers and




Fig. 8. Comparison of a two-axis *Phong* approximation (top) against a ray-traced ground truth solution (bottom).

modern real-time rendering engines make use of physically based BRDFs built upon the microfacet theory. Rough surfaces rendered with a microfacet BRDF exhibit longer specular stretches, more representative of the real phenomenon. The core of the theory relies on the definition of the half vector \mathbf{h} linking the micro geometry variation with the incoming radiance and the viewing direction. Another key aspect is the definition of the *normal distribution function* $D(\mathbf{h})$, responsible for the shape and the brightness of specular highlights. In this section, we demonstrate that microfacet BRDFs can be well approximated using *Irradiance Tensors* theory. Combined with our approximations, we propose a method that can accurately represent the highlight shape, especially the elongated specular stretches viewed at grazing angle, as predicted by the microfacet theory, and at a quality close to the ground truth.

To that end, we consider the axial moment expressed in the half vector space. Following Equation (1), and after proper normalization, this corresponds to the integration of the well known *Blinn-Phong* distribution D_{Blinn}

$$\frac{n+2}{2\pi} M^n(\Omega'_A, \mathbf{n}) = \int_{\Omega'_A} \frac{n+2}{2\pi} (\mathbf{h} \cdot \mathbf{n})^n d\mathbf{h} = \int_{\Omega'_A} D_{\text{Blinn}}(\mathbf{h}) d\mathbf{h}. \quad (23)$$

Given that $d\mathbf{h} = d\mathbf{i}/(4(\mathbf{h} \cdot \mathbf{v}))$, this is equivalent to integrating

$$\int_{\Omega_A} \frac{D_{\text{Blinn}}(\mathbf{h})}{4(\mathbf{h} \cdot \mathbf{v})} d\mathbf{i}. \quad (24)$$

Integrating the axial moment in the half vector space requires the prior knowledge of the transformed spherical region Ω'_A . A naive approach can consist in performing the half vector transform on boundary edge vertices, and evaluate the 1D integral on the newly transformed edges. But as illustrated in Fig. 9, specular highlights get distorted by the

warping distortion introduced by the half vector parameterization. Another possibility is to sample each edge, but it would require a time-consuming per edge evaluation. Previous methods like [16] try to approximate this distortion using anisotropic kernels but it supposes perfect isotropic light emitters only suited for spherical area lights. In our case, the polygonal area lights are not restricted to a specific shape.

8.1 Approximating the Half Vector Warp Distortion

Finding a suitable edge parameterization in half vector space, where axial moment computations can apply, is not straightforward. However, a good approximation can be found. Intuitively, we observe that the distortion reaches its maximum at grazing angles, corresponding to situations where the normal \mathbf{n}_i approaches the surface normal axis \mathbf{n} .

Edge Splitting Strategy. To give the intuition of our method, let consider the great circle gc sustained by a spherical edge and gc' its half vector transformation. If we look at




Fig. 9. Left: The area light's vertices projected in half vector space introduce distortions. Middle: Our edge splitting strategy overcomes the distortion by best approximating the spherical warp for each shaded pixel using only one split. Right: The reference image.

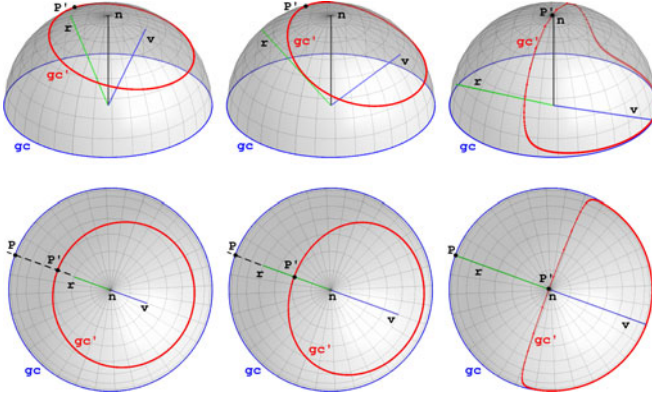


Fig. 10. Illustration of the spherical distortion gc' of the great circle gc produced by the half vector transform. The distortion get its maximum for grazing viewing angles at p' which correspond to the transformation of point p , aligned with the viewing reflection r .

the distortion introduced by the half vector transformation in Fig. 10, we observe that gc' is bent toward the normal axis of gc , with a maximum elevation located at p' , and aligned with the viewing vector v . A simple explanation is that the widest angle spawned by gc with the viewing vector v is found at p . In other words, in the direction of r . This simple observation is the core idea of our edge splitting strategy. Choosing a split position at p will always ensure to get the maximum distortion for an edge in half vector space. The strength of this approach is that a single split is required. Moreover, if the position p is located outside the spherical edge, no split is required and the computational overhead of our solution is greatly reduced. The full edge splitting procedure is described in Algorithm 1.

Algorithm 1. Edge Splitting Procedure

```

for each shading point and each spherical edge  $v_i, v_{i+1}$  do
  Orthogonally project  $r$  to the edge plane with
  normal  $n_i$  at point  $p$ 
  Normalize  $p$ 
  Do the half transform of vertices
   $v_i, v_{i+1} \rightarrow v'_i, v'_{i+1}$ 
  if  $p \in v_i, v_{i+1}$  then
    Split edge at  $p$ 
    Do the half transform  $p \rightarrow p'$ 
    Evaluate edge integral for  $v'_i, p$  and  $p, v'_{i+1}$ 
  else
    */ Do not split */
    Evaluate edge integral  $v'_i, v'_{i+1}$ 

```

8.2 Approximation of Microfacet Specular Distributions

A broad range of microfacet distribution functions found in the literature can be fairly well approximated and integrated by means of axial moment over a spherical region.

Beckmann Approximation. The Beckmann distribution is a peak shape that roughly corresponds to a Blinn-Phong distribution for roughness values $m < 0.5$. A decent integration approximation, using a single axial moment, can be obtained by mapping the Beckmann roughness m to the cosine power exponent n . Noting that $n \approx 2/m^2 - 2$, we obtain

TABLE 3
Timing Overhead per Edge (@720p, Full Coverage + Clipping) for Microfacet Distributions Compared to Phong

Specular distribution	GTX 580	GTX 980Ti
Phong	$\times 1$	$\times 1$
Blinn-Phong	$\times 1.26$	$\times 1.36$
GGX	$\times 2.25$	$\times 1.92$

$$\int_{\Omega'_A} D_{\text{Beckmann}}(\mathbf{h}) d\mathbf{h} \approx \frac{1}{\pi m^2} M^n(\Omega'_A, \mathbf{n}). \quad (25)$$

GGX Approximation. The GGX/Towbridge-Reitz distribution [23] corresponds to an ellipsoid peak shape function producing smoother specular highlights that better match experimental measurements from real materials. At the difference to Blinn-Phong, the distribution has a smoother fall-off, converging to c^2 at the domain bound when $\mathbf{h} \cdot \mathbf{n} = 0$. To mimic this behavior, we split the distribution into one constant term $c_0 = c^2$ and one lobe term corresponding to the GGX distribution shifted down to 0. The integration of the constant term reverts to the calculation of the solid angle sustained by the area-light in the half vector domain. The integration of the second term is approximated by a weighted sum of two axial moments of order n_1 and n_2 where n_2 has a wider eccentricity compared to n_1 to reproduce the GGX smoothness.

$$\int_{\Omega'_A} D_{\text{GGX}}(\mathbf{h}) d\mathbf{h} \approx \frac{1}{\pi} \left(c_0 \Omega'_A + c_1 \frac{n_1 + 2}{2} M^{n_1}(\Omega'_A, \mathbf{n}) + c_2 \frac{n_2 + 2}{2} M^{n_2}(\Omega'_A, \mathbf{n}) \right) \quad (26)$$

The resulting integral is normalized so that $\int D \mathbf{h} \cdot \mathbf{n} = 1$ so the weights c_1 and c_2 are chosen accordingly such as $c_0 + c_1 + c_2 = 1$. Using a least square fitting method, we found that

$$\begin{aligned} n_1 &= \frac{2}{c^2} - 2 & \text{with weights} & \quad c_1 = 0.7 (1 - c^2) \\ n_2 &= n_1/10 & & \quad c_2 = 0.3 (1 - c^2), \end{aligned}$$

provide a decent approximation whatever the eccentricity parameter c .

Modeling More Complex Distributions. Integrating directly in the half-vector space allows for the shading of microfacet surfaces with anisotropic properties. Such anisotropic distributions [24] [25] can be obtained by scaling/stretching the polygon along desired axes. Modeling more complex distribution in that space can be achieved numerically by best-fitting LTC matrices from [4] or analytically by combining several axes as we did in Section 7.1.

8.3 Results and Improvements

We implemented and tested our solution on a GPU NVIDIA GTX 580 and a GTX 980Ti. The Table 3 gives the timing overhead compared to Phong distribution. In contrast to Phong, only one horizon clipping is performed around \mathbf{n} . As a result, combined with our low-cost edge splitting approach, our solution has a limited computational overhead.




Fig. 11. Surface lit with a Blinn-Phong and GGX microfacet distribution rendered using our analytic approximations compared against a reference ray-traced solution.

We also compared our solution with reference images obtained with a ray-traced solution. As shown in Fig. 11, the elongated specular stretches predicted by the microfacet theory are faithfully reproduced with an accuracy close to the reference in most situation. However, under certain roughness and geometric configurations we can notice a lack of brightness, especially when the viewing reflection is close to an edge border and when the surface is rough (see Fig. 13). This phenomenon tends to disappear as the roughness goes to 0 ($n \rightarrow \infty$). These artifacts result from an underestimation of the theoretical solid angle in the half-vector space. As shown in Fig. 12, our edge-splitting strategy fails to capture important distortion of the great circle sustained by an edge.

Edge-Split Balancing Strategy. To reduce the visual artifacts, we propose a simple balancing strategy that captures the inflection points on the half-transformed great circle. The idea is to balance the split position between the mid point \mathbf{p}_m of v_i, v_{i+1} and the split position \mathbf{p} prior to the half-transform. The balancing heuristic depends on the surface roughness m and the angle between \mathbf{v} and the edge plane with normal \mathbf{t}_i . The balancing strategy is summarized in the Algorithm 2.

Algorithm 2. Edge-Split Balancing Procedure

$m = \sqrt{2.0/(n+2)}$
 $k = 2 |\mathbf{t}_i \cdot \mathbf{v}|$
 $\mathbf{p}_m = (v_i + v_{i+1}) / \|v_i + v_{i+1}\|$
 $\mathbf{p} = (1 - k m) \mathbf{p} + k m \mathbf{p}_m$
 $\mathbf{p} = \mathbf{p} / \|\mathbf{p}\|$

As illustrated in Fig. 13, the balancing strategy overcomes most of the visual artifacts. The main advantage is that no additional split is required. However, in some extreme cases, the single split approach still shows some differences at extreme grazing angles, or when the light source is very large.

9 SPHERICAL AND DISC AREA LIGHTS SUPPORT

The analytic approximations we have developed so far are restricted to polygonal light emitters. Light emitters based

on analytic shapes such as *Sphere* and *Disc*, represent an interesting class of area-lights that extends the set of luminaries representation. Note that Arvo [21] proposed another closed form expression to evaluate the axial moment over a spherical light. However, the solution is based on another parameterization and recurrence evaluation solved in $O(n)$ time. Furthermore, the method cannot handle non-uniform scaling operation.

We propose a simple method to shade specular surfaces by spherical and disc area lights. Our method leverages the polygonal approach with our approximations, supporting the partial visibility of the lights and non-uniform scaling operations.

9.1 Spinning Polygon Strategy

Our approach is inspired by optical illusions produced by high-speed spinning rotations. The idea, depicted in Fig. 14, is to give the illusion of a sphere or a disc by considering the spinning of a k -sided polygon around a unit disc normal axis and evaluate the axial moment over the resulting polygon. The orientation θ_m of the polygon is computed at each shading point and should theoretically be chosen such as

$$\theta_m = \arg \max M^n(P(\theta), \mathbf{r}).$$




Fig. 12. Solid projection of a rectangular area light on the hemisphere (in blue) and its corresponding theoretical projection in the half-vector space (in red). On the left, the edge split approximation (in green) under-estimates the theoretical projection when the surface is rough and when \mathbf{r} is at grazing angle against the edge plane. Using a simple balancing heuristic, controlled by \mathbf{r} (middle) and by the roughness (right), we better recover the theoretical solid angle of the half-vector transformed polygon.




Fig. 13. Under certain roughness and geometric configurations the edge splitting strategy may exhibits some visual artifacts (split). Thanks to a simple balancing heuristic (split+bal) those artifacts are greatly reduced (left+middle). However, in some extreme configurations, some visual differences are still perceptible (right).

Disc Area Light. Determining the best orientation θ_m is a nontrivial maximization problem that depends on several parameters difficult to solve in real-time. However, according to experimental measurements, we noticed that when the roughness goes to 0 ($n \rightarrow \infty$), the best orientation is found toward r' , the intersection of the (p, r) line with the disc plane. Conversely, when the roughness goes to 1 ($n \rightarrow 1$) the best orientation is found towards p' the perpendicular projection of p on the disc plane. By setting-up a linear blending between the two positions driven by the roughness m we can obtain a decent approximation for the best orientation. Another problem is the area difference between the polygon and the disc resulting in an underestimation of the brightness when the surface is rough. Windowed by the *Cosine lobe* eccentricity, the difference is reduced when the surface is highly specular. To reduce the difference in all scenarios, we linearly scale the polygon such as $\text{area}(P) = \text{area}(D)$ as the roughness goes to 0.

Sphere Area Light. For a sphere light, the procedure is roughly the same as for the disc. The main difference lies in the orientation of the disc, facing the shading point p . Also, a scaling factor s is applied to the unit disc to take into account of the solid angle sustained by the unit sphere (see Fig. 15). For both luminaries, non uniform scaling operations are simply supported by transforming p and r into the area light local space prior the orientation estimation. The resulting polygon is then back-transformed into the world space. The full procedure for evaluating the shading from disc and sphere area lights is described in Algorithm 3.

Algorithm 3. Disc and Sphere Light Shading

```

 $m = \sqrt{(2/(n+2))}$ 
 $s = 1$  /* default disc scale */
for each shading point  $p$  do
    /* handle non uniform-scaling */
    Transform  $p$  and  $r$  in area light local space
    if Sphere then
         $n_D = \text{normalize}(p)$ 
         $d = \text{length}(p)$ 
         $s = d/\sqrt{d^2 - 1}$ 
         $k = |n_D \cdot p|/d$  /* prevent early clip */
         $s = s(1 - m k + m k \sqrt{\pi/2})$  /* +area diff */
         $r' = \text{intersection with the disc plane in the } r \text{ direction}$ 
         $p' = \text{orthogonal projection of } p \text{ to the disc plane}$ 
         $v_0 = s \text{ normalize}((m-1)r' + m p')$ 
         $v_1 = v_0 \times n_D$ 
         $v_2 = -v_0$ 
         $v_3 = -v_1$ 
        Back-transform  $v_0 \rightarrow v_3$  in world space
        Compute the axial moment on polygon  $v_0 \rightarrow v_3$ 

```

9.2 Results and Limitations

Our spinning approach (see Fig. 1d) provides convincing specular highlights for spherical and disc area light sources with $k = 4$. The partial visibility is properly handled and non-uniform scaling operations allows the representation of ellipses and ellipsoid shaped area lights. In terms of performance, we didn't notice a significant difference compared to the quad evaluation. However, our approach has some limitations. First, our method to compensate for




Fig. 14. Description of our spinning approach. We give illusion of a disc by considering the spinning of a k -sided polygon P around the disc normal n_d . The orientation of P (here a quad) in the unit disc is chosen such as $\arg \max_{\theta} M^n(P(\theta), r)$. The shading at point p is then estimated by evaluating the axial moment on the resulting polygon P .




Fig. 15. The axial moment over a spherical light is estimated by facing the unit disc toward the shading point p (left). The disc (and hence the polygon P) is then scaled by s to take into account of the solid angle of the sphere (right).




Fig. 16. Disc (top) and sphere (bottom) area lights simulated using our spinning quad technique. Our method produces convincing results compared to ground truth ray-traced references (here with a *Phong* distribution), supports non-uniform scale operations and horizon clipping.

the area difference is just an approximation. As we do not integrate the circular shape we observe slight brightness differences especially with moderate glossy surfaces. Second, when the predicted position to orient the polygon is too close to the disc center, the excessive rotations of the polygon on close pixels results in perceptible brightness variations. With a *Phong* distribution, this variation is mostly perceptible when the disc highlight is viewed at grazing angle. With microfacet distributions, the phenomenon is highly perceptible due to the distortion introduced by the half vector transform. This variation disappears as n goes to infinity.

10 CONCLUSION & FUTURE WORK

We presented efficient and accurate analytic approximations for the surface shading from polygonal light sources. Our method is flexible and fast enough for high quality demanding real-time applications. In particular, we showed that the edge integrals of Arvo can be accurately approximated and evaluated in constant-time and that the integration of multiple axis-oriented lobes can be easily approximated using derivation from SG operators. We also demonstrated that the *Phong* restriction can be overcome by approximating the half vector warp distortions using a single edge-split strategy. However, in some extreme configurations, the edge-splitting strategy may exhibit undesirable artifacts. To overcome these defects, we plan to explore alternative parameterization of the

great circle in the half vector space. Our goal is to avoid the splitting and better match the distortions introduced by the spherical warp. For spherical and disc area light sources, our spinning quads produce convincing results for *Phong* distribution at roughly the same cost as the quad. But in some configurations, the excessive rotation of the quad introduces undesirable brightness variations we need to address. We found for instance that finding the optimal orientation can be reduced to a 1D problem by just considering the integration along the diagonal of the quad.

Other challenges still remain that would be worth exploring in the future. First, soft shadows are ignored with our method. One solution would be to back-project the scene geometry onto the area light and perform a negative



Fig. 17. Limitations of our polygon spinning strategy. In some particular cases, the excessive rotation of the polygon in a close area introduces undesirable brightness variations. With *Phong* distribution (left), the phenomenon is perceptible only with discs when viewed from grazing angle. With microfacet distributions (middle + right), the phenomenon is perceptible in many configuration.

contour integration along the geometry silhouette. Textured area lights is also a hard problem for which no satisfying solution exists yet. One possibility with our approach is to modulate the specular term with pre-integrated mip-mapped textures as done in [1] and [4]. One other approach would be to look for the varying luminaries derivations introduced by Arvo [21] and developed by Chen and Arvo [26]. Finally, some broader lighting problems such as real-time *environment lighting* or interactive *Global Illumination* would be interesting to address. We believe that our approximation framework can be particularly well adapted to these techniques and may overcome some of the issues encountered with *Spherical Gaussians* or *VPLs* approaches.

REFERENCES

- [1] M. Drobot, "Physically based area lights," in *Proc. GPU Pro 5*, 2014, pp. 67–100.
- [2] B. Karis, "Real shading in unreal engine 4," in *Proc. ACM SIGGRAPH Courses*, 2013, pp. 22:1–22:8.
- [3] J. Arvo, "Applications of irradiance tensors to the simulation of non-lambertian phenomena," in *Proc. 22nd Annu. Conf. Comput. Graph. Interactive Techn.*, 1995, pp. 335–342.
- [4] E. Heitz, J. Dupuy, S. Hill, and D. Neubelt, "Real-time polygonal-light shading with linearly transformed cosines," *ACM Trans. Graph.*, vol. 35, no. 4, pp. 41:1–41:8, Jul. 2016.
- [5] J. Arvo, "Stratified sampling of spherical triangles," in *Proc. 22nd Annu. Conf. Comput. Graph. Interactive Techn.*, 1995, pp. 437–438.
- [6] C. Ureña, M. Fajardo, and A. King, "An area-preserving parametrization for spherical rectangles," *Comput. Graph. Forum*, vol. 32, no. 4, pp. 59–66, 2013.
- [7] A. Keller, "Instant radiosity," in *Proc. 24th Annu. Conf. Comput. Graph. Interactive Techn.*, 1997, pp. 49–56.
- [8] B. Walter, S. Fernandez, A. Arbree, K. Bala, M. Donikian, and D. P. Greenberg, "Lightcuts: A scalable approach to illumination," in *Proc. Annu. Conf. Comput. Graph. Interactive Techn.*, 2005, pp. 1098–1107.
- [9] G. Nichols and C. Wyman, "Direct illumination from dynamic area lights," in *Proc. Annu. Conf. Comput. Graph. Interactive Techn.*, 2009, pp. 82:1–82:1.
- [10] T. Ritschel, T. Grosch, J. Kautz, and H.-P. Seidel, "Interactive global illumination based on coherent surface shadow maps," in *Proc. Graph. Interface*, 2008, pp. 185–192.
- [11] K. Picott, "Extensions of the linear and area lighting models," *IEEE Comput. Graph. Appl.*, vol. 12, no. 2, pp. 31–38, Mar. 1992.
- [12] H. Bao and Q. Peng, "Shading models for linear and area light sources," *Comput. Graph.*, vol. 17, no. 2, pp. 137–145, 1993.
- [13] T. Tanaka and T. Takahashi, "Fast analytic shading and shadowing for area light sources," *Comput. Graph. Forum*, vol. 16, pp. C231–C240, 1997.
- [14] P. Poulin and J. Amanatides, "Shading and shadowing with linear light sources," *Comput. Graph.*, vol. 15, no. 2, pp. 259–265, 1991.
- [15] J. M. Snyder, "Area light sources for real-time graphics," Microsoft Research, Redmond, WA, USA, Tech. Rep. MSR-TR-96-11, 1996.
- [16] J. Wang, P. Ren, M. Gong, J. Snyder, and B. Guo, "All-frequency rendering of dynamic, spatially-varying reflectance," in *Proc. SIGGRAPH Asia*, 2009, pp. 133:1–133:10.
- [17] K. Xu, W.-L. Sun, Z. Dong, D.-Y. Zhao, R.-D. Wu, and S.-M. Hu, "Anisotropic spherical gaussians," *ACM Trans. Graph.*, vol. 32, no. 6, pp. 209:1–209:11, Nov. 2013.
- [18] Y. Tokuyoshi, "Virtual spherical gaussian lights for real-time glossy indirect illumination," in *Proc. SIGGRAPH Asia Techn. Briefs*, 2014, pp. 17:1–17:4.
- [19] R. Wang, et al., "Analytic double product integrals for all-frequency relighting," *IEEE Trans. Vis. Comput. Graph.*, vol. 19, no. 7, pp. 1133–1142, Jul. 2013.
- [20] K. Xu, Y.-P. Cao, L.-Q. Ma, Z. Dong, R. Wang, and S.-M. Hu, "A practical algorithm for rendering interreflections with all-frequency BRDFs," *ACM Trans. Graph.*, vol. 33, no. 1, Feb. 2014, Art. no. 10.
- [21] J. Arvo, "Analytic methods for simulated light transport," PhD dissertation, Yale Univ., New Haven, CT, USA, 1995.
- [22] S. Meunier, R. Perrot, L. Aveneau, D. Meneveaux, and D. Ghazanfarpour, "Technical section: Cosine lobes for interactive direct lighting in dynamic scenes," *Comput. Graph.*, vol. 34, no. 6, pp. 767–778, Dec. 2010.
- [23] B. Walter, S. R. Marschner, H. Li, and K. E. Torrance, "Microfacet models for refraction through rough surfaces," in *Proc. 18th Eurographics Conf. Rendering Techn.*, 2007, pp. 195–206.
- [24] M. Ashikhmin and P. Shirley, "An anisotropic phong BRDF model," *J. Graph. Tools*, vol. 5, no. 2, pp. 25–32, Feb. 2000.
- [25] E. P. F. Lafortune, S.-C. Foo, K. E. Torrance, and D. P. Greenberg, "Non-linear approximation of reflectance functions," in *Proc. 24th Annu. Conf. Comput. Graph. Interactive Techn.*, 1997, pp. 117–126.
- [26] M. Chen and J. Arvo, "A closed-form solution for the irradiance due to linearly-varying luminaires," in *Proc. Eurographics Workshop Rendering Techn.*, 2000, pp. 137–148.



Pascal Lecocq received the PhD degree in computer science from the University of Paris-Est Marne-La-Valle, France, in 2001. He is a senior scientist with Technicolor Research & Innovation. His main activity at Technicolor is to study and develop interactive previs solutions for VFX and animation productions. His research is focused on lighting, high-quality shadows, and participating media rendering.



Arthur Dufay received the master's degree in computer-science from the University of Lyon, France. He is working toward the PhD degree at Technicolor Research & Innovation, since June 2014. His research is focused on real-time rendering and GPU acceleration techniques for global-illumination problematics for previs and production rendering.



Gaël Sourimant received the PhD degree in computer vision from the University of Rennes, France, in 2007. He is a senior engineer with Technicolor Research & Innovation. His work is focused on high-quality real-time rendering techniques for previs and VR applications.



Jean-Eudes Marvie received the PhD degree from the University of Rennes, in 2004. He is a principal scientist with Technicolor Research & Innovation. His research interests include real time rendering, large models visualization and generation, procedural modeling and rendering, and virtual reality. He is currently leading a team of 10 researchers and engineers, applying these techniques to the field of interactive pre-visualization for cinema production and home VR experiences.

► For more information on this or any other computing topic, please visit our Digital Library at www.computer.org/publications/dlib.